Newton's Method

Ryan Tibshirani
Convex Optimization 10-725

Last time: dual correspondences

Given a function $f : \mathbb{R}^n \to \mathbb{R}$, we define its conjugate $f^* : \mathbb{R}^n \to \mathbb{R}$,

$$f^*(y) = \max_x \; y^T x - f(x)$$

Properties and examples:

- Conjugate f^* is always convex (regardless of convexity of f)
- When f is a quadratic in $Q \succ 0$, f^* is a quadratic in Q^{-1}
- When f is a norm, f^* is indicator of the dual norm unit ball
- When f is closed and convex, $x \in \partial f^*(y) \iff y \in \partial f(x)$

Relationship to duality (also called Fenchel duality):

$$\text{Primal}: \quad \min_x \; f(x) + g(x)$$
$$\text{Dual}: \quad \max_u \; -f^*(u) - g^*(-u)$$

Newton's method

Given unconstrained, smooth convex optimization

$$\min_x \; f(x)$$

where f is convex, twice differentiable, and $\mathrm{dom}(f) = \mathbb{R}^n$. Recall that gradient descent chooses initial $x^{(0)} \in \mathbb{R}^n$, and repeats

$$x^{(k)} = x^{(k-1)} - t_k \cdot \nabla f(x^{(k-1)}), \quad k = 1, 2, 3, \ldots$$

In comparison, Newton's method repeats

$$x^{(k)} = x^{(k-1)} - \left(\nabla^2 f(x^{(k-1)})\right)^{-1} \nabla f(x^{(k-1)}), \quad k = 1, 2, 3, \ldots$$

Here $\nabla^2 f(x^{(k-1)})$ is the Hessian matrix of f at $x^{(k-1)}$

Newton's method interpretation

Recall the motivation for gradient descent step at x: we minimize the quadratic approximation

$$f(y) \approx f(x) + \nabla f(x)^T (y - x) + \frac{1}{2t} \|y - x\|_2^2$$

over y, and this yields the update $x^+ = x - t \nabla f(x)$

Newton's method uses in a sense a better quadratic approximation

$$f(y) \approx f(x) + \nabla f(x)^T (y - x) + \frac{1}{2} (y - x)^T \nabla^2 f(x)(y - x)$$

and minimizes over y to yield $x^+ = x - (\nabla^2 f(x))^{-1} \nabla f(x)$

Consider minimizing $f(x) = (10x_1^2 + x_2^2)/2 + 5\log(1 + e^{-x_1-x_2})$
(this must be a nonquadratic ... why?)

We compare gradient descent (black) to Newton's method (blue), where both take steps of roughly same length

Outline

Today:

- Interpretations and properties
- Backtracking line search
- Convergence analysis
- Equality-constrained Newton
- Quasi-Newton preview

Linearized optimality condition

Aternative interpretation of Newton step at x: we seek a direction v so that $\nabla f(x + v) = 0$. Let $F(x) = \nabla f(x)$. Consider linearizing F around x, via first-order approximation

$$0 = F(x + v) \approx F(x) + DF(x)v$$

Solving for v yields $v = -(DF(x))^{-1}F(x) = -(\nabla^2 f(x))^{-1}\nabla f(x)$

(From B & V page 486)

History: work of Newton (1685) and Raphson (1690) originally focused on finding roots of polynomials. Simpson (1740) applied this idea to general nonlinear equations, and minimization by setting the gradient to zero

Affine invariance of Newton's method

Important property Newton's method: affine invariance. Given f, nonsingular $A \in \mathbb{R}^{n \times n}$. Let $x = Ay$, and $g(y) = f(Ay)$. Newton steps on g are

$$
\begin{aligned}
y^+ &= y - \left(\nabla^2 g(y)\right)^{-1} \nabla g(y) \\
&= y - \left(A^T \nabla^2 f(Ay) A\right)^{-1} A^T \nabla f(Ay) \\
&= y - A^{-1} \left(\nabla^2 f(Ay)\right)^{-1} \nabla f(Ay)
\end{aligned}
$$

Hence

$$
Ay^+ = Ay - \left(\nabla^2 f(Ay)\right)^{-1} \nabla f(Ay)
$$

i.e.,

$$
x^+ = x - \left(\nabla^2 f(x)\right)^{-1} \nabla f(x)
$$

So progress is independent of problem scaling. This is not true of gradient descent!

Newton decrement

At a point x, we define the Newton decrement as

$$\lambda(x) = \left(\nabla f(x)^T \big(\nabla^2 f(x) \big)^{-1} \nabla f(x) \right)^{1/2}$$

This relates to the difference between $f(x)$ and the minimum of its quadratic approximation:

$$
\begin{aligned}
f(x) - \min_y & \left(f(x) + \nabla f(x)^T (y - x) + \frac{1}{2}(y - x)^T \nabla^2 f(x)(y - x) \right) \\
&= f(x) - \left(f(x) - \frac{1}{2}\nabla f(x)^T \big(\nabla^2 f(x) \big)^{-1} \nabla f(x) \right) \\
&= \frac{1}{2}\lambda(x)^2
\end{aligned}
$$

Therefore can think of $\lambda^2(x)/2$ as an approximate upper bound on the suboptimality gap $f(x) - f^\star$

Another interpretation of Newton decrement: if Newton direction is $v = -(\nabla^2 f(x))^{-1} \nabla f(x)$, then

$$\lambda(x) = \left(v^T \nabla^2 f(x) v\right)^{1/2} = \|v\|_{\nabla^2 f(x)}$$

i.e., $\lambda(x)$ is the length of the Newton step in the norm defined by the Hessian $\nabla^2 f(x)$

Note that the Newton decrement, like the Newton steps, are affine invariant; i.e., if we defined $g(y) = f(Ay)$ for nonsingular A, then $\lambda_g(y)$ would match $\lambda_f(x)$ at $x = Ay$

Backtracking line search

So far we've seen pure Newton's method. This need not converge. In practice, we use damped Newton's method (typically just called Newton's method), which repeats

$$x^+ = x - t\big(\nabla^2 f(x)\big)^{-1}\nabla f(x)$$

Note that the pure method uses $t = 1$

Step sizes here are chosen by backtracking search, with parameters $0 < \alpha \le 1/2$, $0 < \beta < 1$. At each iteration, start with $t = 1$, while

$$f(x + tv) > f(x) + \alpha t\nabla f(x)^T v$$

we shrink $t = \beta t$, else we perform the Newton update. Note that here $v = -(\nabla^2 f(x))^{-1}\nabla f(x)$, so $\nabla f(x)^T v = -\lambda^2(x)$

Example: logistic regression

Logistic regression example, with $n = 500$, $p = 100$: we compare gradient descent and Newton's method, both with backtracking

Newton's method: in a totally different regime of convergence...!

Convergence analysis

Assume that f convex, twice differentiable, having $\mathrm{dom}(f) = \mathbb{R}^n$, and additionally

- ∇f is Lipschitz with parameter L
- f is strongly convex with parameter m
- $\nabla^2 f$ is Lipschitz with parameter M

Theorem: Newton's method with backtracking line search satisfies the following two-stage convergence bounds

$$f(x^{(k)}) - f^\star \leq \begin{cases} (f(x^{(0)}) - f^\star) - \gamma k & \text{if } k \leq k_0 \\ \dfrac{2m^3}{M^2}\left(\dfrac{1}{2}\right)^{2^{k-k_0+1}} & \text{if } k > k_0 \end{cases}$$

Here $\gamma = \alpha\beta^2\eta^2 m/L^2$, $\eta = \min\{1, 3(1-2\alpha)\}m^2/M$, and k_0 is the number of steps until $\|\nabla f(x^{(k_0+1)})\|_2 < \eta$

In more detail, convergence analysis reveals $\gamma > 0$, $0 < \eta \leq m^2/M$ such that convergence follows two stages

- Damped phase: $\|\nabla f(x^{(k)})\|_2 \geq \eta$, and

$$f(x^{(k+1)}) - f(x^{(k)}) \leq -\gamma$$

- Pure phase: $\|\nabla f(x^{(k)})\|_2 < \eta$, backtracking selects $t = 1$, and

$$\frac{M}{2m^2} \|\nabla f(x^{(k+1)})\|_2 \leq \left(\frac{M}{2m^2} \|\nabla f(x^{(k)})\|_2 \right)^2$$

Note that once we enter pure phase, we won't leave, because

$$\frac{2m^2}{M} \left(\frac{M}{2m^2} \eta \right)^2 \leq \eta$$

when $\eta \leq m^2/M$

Unraveling this result, what does it say? To get $f(x^{(k)}) - f^\star \leq \epsilon$, we need at most

$$\frac{f(x^{(0)}) - f^\star}{\gamma} + \log\log(\epsilon_0/\epsilon)$$

iterations, where $\epsilon_0 = 2m^3/M^2$

- This is called quadratic convergence. Compare this to linear convergence (which, recall, is what gradient descent achieves under strong convexity)

- The above result is a local convergence rate, i.e., we are only guaranteed quadratic convergence after some number of steps k_0, where $k_0 \leq \frac{f(x^{(0)}) - f^\star}{\gamma}$

- Somewhat bothersome may be the fact that the above bound depends on L, m, M, and yet the algorithm itself does not ...

Self-concordance

A scale-free analysis is possible for self-concordant functions: on $\mathbb{R}$, a convex function f is called self-concordant if

$$|f'''(x)| \le 2f''(x)^{3/2} \quad \text{for all } x$$

and on $\mathbb{R}^n$ is called self-concordant if its projection onto every line segment is so

Theorem (Nesterov and Nemirovskii): Newton's method with backtracking line search requires at most

$$C(\alpha, \beta)\big(f(x^{(0)}) - f^\star\big) + \log\log(1/\epsilon)$$

iterations to reach $f(x^{(k)}) - f^\star \le \epsilon$, where $C(\alpha, \beta)$ is a constant that only depends on α, β

What kind of functions are self-concordant?

- Linear and quadratic functions
- $f(x) = -\sum_{i=1}^{n} \log(x_i)$ on $\mathbb{R}_{++}^n$
- $f(X) = -\log(\det(X))$ on $\mathbb{S}_{++}^n$
- If g is self-concordant, then so is $f(x) = g(Ax + b)$
- In the definition of self-concordance, we can replace factor of 2 by a general $\kappa > 0$
- If g is κ-self-concordant, then we can rescale: $f(x) = \frac{\kappa^2}{4} g(x)$ is self-concordant (2-self-concordant)

Comparison to first-order methods

At a high-level:

- **Memory**: each iteration of Newton's method requires $O(n^2)$ storage ($n \times n$ Hessian); each gradient iteration requires $O(n)$ storage (n-dimensional gradient)

- **Computation**: each Newton iteration requires $O(n^3)$ flops (solving a dense $n \times n$ linear system); each gradient iteration requires $O(n)$ flops (scaling/adding n-dimensional vectors)

- **Backtracking**: backtracking line search has roughly the same cost, both use $O(n)$ flops per inner backtracking step

- **Conditioning**: Newton's method is not affected by a problem's conditioning, but gradient descent can seriously degrade

Back to logistic regression example: now x-axis is parametrized in terms of time taken per iteration

Each gradient descent step is $O(p)$, but each Newton step is $O(p^3)$

Sparse, structured problems

When the inner linear systems (in Hessian) can be solved efficiently and reliably, Newton's method can strive

For example, if $\nabla^2 f(x)$ is sparse/structured for all x, say banded, then both memory and computation are $O(n)$ per Newton iteration

What functions admit a structured Hessian? Two examples:

- If $g(\beta) = f(X\beta)$, then $\nabla^2 g(\beta) = X^T \nabla^2 f(X\beta) X$. Hence if X is a structured predictor matrix and $\nabla^2 f$ is diagonal, then $\nabla^2 g$ is structured

- If we seek to minimize $f(\beta) + g(D\beta)$, where $\nabla^2 f$ is diagonal, g is not smooth, and D is a structured penalty matrix, then the Lagrange dual function is $-f^*(-D^T u) - g^*(-u)$. Often $\nabla^2 f^*$ will be diagonal (e.g., when $f(\beta) = \sum_{i=1}^p f_i(\beta_i)$) so the Hessian in dual will be structured

Equality-constrained Newton's method

Consider now a problem with equality constraints, as in

$$\min_x \ f(x) \quad \text{subject to} \quad Ax = b$$

Several options:

- **Eliminating equality constraints**: write $x = Fy + x_0$, where F spans null space of A, and $Ax_0 = b$. Solve in terms of y

- **Deriving the dual**: can check that the Lagrange dual function is $-f^*(-A^T v) - b^T v$, and strong duality holds. With luck, we can express $x^\star$ in terms of $v^\star$

- **Equality-constrained Newton**: in many cases, this is the most straightforward option

In equality-constrained Newton's method, we start with $x^{(0)}$ such that $Ax^{(0)} = b$. Then we repeat the updates

$$x^+ = x + tv, \quad \text{where}$$

$$v = \operatorname*{argmin}_{Az=0} \ \nabla f(x)^T(z - x) + \frac{1}{2}(z - x)^T \nabla^2 f(x)(z - x)$$

This keeps x^+ in feasible set, since $Ax^+ = Ax + tAv = b + 0 = b$

Furthermore, v is the solution to minimizing a quadratic subject to equality constraints. We know from KKT conditions that v satisfies

$$\left[\begin{array}{cc} \nabla^2 f(x) & A^T \\ A & 0 \end{array} \right] \left[\begin{array}{c} v \\ w \end{array} \right] = \left[\begin{array}{c} -\nabla f(x) \\ 0 \end{array} \right]$$

for some w. Hence Newton direction v is again given by solving a linear system in the Hessian (albeit a bigger one)

Quasi-Newton methods

If the Hessian is too expensive (or singular), then a quasi-Newton method can be used to approximate $\nabla^2 f(x)$ with $H \succ 0$, and we update according to

$$x^+ = x - tH^{-1}\nabla f(x)$$

- Approximate Hessian H is recomputed at each step. Goal is to make H^{-1} cheap to apply (possibly, cheap storage too)

- Convergence is fast: superlinear, but not the same as Newton. Roughly n steps of quasi-Newton make same progress as one Newton step

- Very wide variety of quasi-Newton methods; common theme is to "propogate" computation of H across iterations

References and further reading

- S. Boyd and L. Vandenberghe (2004), "Convex optimization", Chapters 9 and 10
- Y. Nesterov (1998), "Introductory lectures on convex optimization: a basic course", Chapter 2
- Y. Nesterov and A. Nemirovskii (1994), "Interior-point polynomial methods in convex programming", Chapter 2
- J. Nocedal and S. Wright (2006), "Numerical optimization", Chapters 6 and 7
- L. Vandenberghe, Lecture notes for EE 236C, UCLA, Spring 2011-2012